


Actionable Uncertainty for LLM-based Recommendation

Controlled same-candidate ranking and evidence gates



Claim boundary: C-CRP tests whether task-grounded calibrated uncertainty improves controlled candidate ranking reliability.

The pipeline does not assert full-catalog recommendation, generative title recommendation, or backbone novelty.